

✓ Thyroid Disease Classification using Neural Network

Data Loading and Preprocessing

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
from sklearn.utils import resample

import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense

# Load the dataset
data_path = 'sample_data/thyroidDF.csv'
df = pd.read_csv(data_path)

# Print the initial shape of the dataset
print("Initial dataset shape:", df.shape)

Initial dataset shape: (9172, 31)
```

```
print(df.iloc[:, :4], df['target'])
```

	age	sex	on_thyroxine	query_on_thyroxine
0	29	F	f	f
1	29	F	f	f
2	41	F	f	f
3	36	F	f	f
4	32	F	f	f
...
9167	56	M	f	f
9168	22	M	f	f
9169	69	M	f	f
9170	47	F	f	f
9171	31	M	f	f

```
[9172 rows x 4 columns] 0      -
1      -
2      -
3      -
4      S
...
9167   -
9168   -
9169   I
9170   -
9171   -
Name: target, Length: 9172, dtype: object
```

```
# Handling missing values
for column in df.columns:
    if df[column].dtype in ['float64', 'int64']:
        df[column].fillna(df[column].median(), inplace=True)
    elif df[column].dtype == 'object':
        df[column].fillna(df[column].mode()[0], inplace=True)
```

```
# Encode categorical variables using Label Encoder
for column in df.columns:
    if df[column].dtype == 'object':
        encoder = LabelEncoder()
        df[column] = encoder.fit_transform(df[column])
```

```
# Balancing skewed dataset
# Separate majority and minority classes
df_majority = df[df.target == df.target.mode()[0]]
df_minority = df[df.target != df.target.mode()[0]]

# Upsample minority class
df_minority_upsampled = resample(df_minority,
                                replace=True,      # sample with replacement
                                n_samples=len(df_majority),  # to match majority
                                random_state=123) # reproducible results

# Combine majority class with upsampled minority class
df = pd.concat([df_majority, df_minority_upsampled])

# Shuffle the dataset to avoid any order bias
df = df.sample(frac=1, random_state=42).reset_index(drop=True)

print("New dataset shape:", df.shape) # Print the new shape of the dataset
```

New dataset shape: (13542, 31)

```
# Separate features and target variable
features = df.drop(['target', 'patient_id'], axis=1)
labels = df['target']

# Scale the features
scaler = StandardScaler()
features_scaled = scaler.fit_transform(features)

# Split the data into training and testing sets
trainX, testX, trainY, testY = train_test_split(features_scaled, labels, test_s
```

Model

```
# Define the neural network
model = Sequential([
    Dense(128, activation='relu', input_shape=(trainX.shape[1],)),
    Dense(128, activation='relu'),
    Dense(64, activation='relu'),
    Dense(len(labels.unique()), activation='softmax')
])


model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

# Train the model
history = model.fit(trainX, trainY, epochs=10, validation_data=(testX, testY))


Epoch 1/10
339/339 [=====] - 6s 8ms/step - loss: 1.1593 - acc
Epoch 2/10
339/339 [=====] - 1s 3ms/step - loss: 0.5526 - acc
Epoch 3/10
339/339 [=====] - 1s 4ms/step - loss: 0.4047 - acc
Epoch 4/10
339/339 [=====] - 1s 3ms/step - loss: 0.3284 - acc
Epoch 5/10
339/339 [=====] - 1s 3ms/step - loss: 0.2832 - acc
Epoch 6/10
339/339 [=====] - 1s 3ms/step - loss: 0.2563 - acc
Epoch 7/10
339/339 [=====] - 1s 3ms/step - loss: 0.2363 - acc
Epoch 8/10
339/339 [=====] - 1s 3ms/step - loss: 0.2157 - acc
Epoch 9/10
339/339 [=====] - 1s 4ms/step - loss: 0.1993 - acc
Epoch 10/10
339/339 [=====] - 2s 5ms/step - loss: 0.1815 - acc
```

Plotting Training and Validation Metrics

```
# Plot training and validation loss over epochs
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Training and Validation Loss')
plt.legend()
plt.show()
```



```
# Plot training and validation accuracy over epochs
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.title('Training and Validation Accuracy')
plt.legend()
plt.show()
```



```

predictions = model.predict(testX)
predicted_classes = np.argmax(predictions, axis=1)

accuracy = accuracy_score(testY, predicted_classes)
precision = precision_score(testY, predicted_classes, average='weighted')
recall = recall_score(testY, predicted_classes, average='weighted')
f1 = f1_score(testY, predicted_classes, average='weighted')

print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("F1-score:", f1)

```

```

85/85 [=====] - 0s 2ms/step
Accuracy: 0.9280177187153932
Precision: 0.9324928247906574
Recall: 0.9280177187153932
F1-score: 0.9274882283493258

```

```
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 128)	3840
dense_1 (Dense)	(None, 128)	16512
dense_2 (Dense)	(None, 64)	8256
dense_3 (Dense)	(None, 32)	2080

```

=====
Total params: 30688 (119.88 KB)
Trainable params: 30688 (119.88 KB)
Non-trainable params: 0 (0.00 Byte)

```

